



1.7B

**Best@8**

8B





Figure 5. Validation performance of Qwen3-1.7B-Base (top) and Qwen3-8B-Base (bottom) trained with GRPO and GRPO-DBB. The left plots show Mean@8, and the right plots show Best@8 on the validation set.